

VARIA: Benchmarking Frontier LLMs on Construct-Equivalent Assessment Variant Generation

Dr. Ernesto Lee Miami Dade College elee@mdc.edu

Abstract

In plain terms: if we are going to let AI write personalized versions of tests for every student, those versions have to be more than surface paraphrase — they have to look genuinely different on the outside while still measuring the same underlying skill on the inside. This paper asks whether today’s frontier AI models can actually do that, and offers a way to measure it.

AI-Integrated Authentic Assessment (AIAA) frameworks propose to replace surveillance-based proctoring with per-student task variation: each examinee receives a unique-yet-equivalent performance task, making answer copying useless without invasive monitoring. The framework’s no-surveillance integrity guarantee is conditional on a presumed but unverified capability — that large language models can generate parallel variants of an assessment blueprint that are simultaneously surface-distinct, construct-equivalent, rubric-applicable, and difficulty-matched. We introduce VARIA, a benchmark that operationalises this presumption as four measurable properties of a generated variant set. VARIA evaluates three frontier model families (Anthropic Claude Opus 4.7, OpenAI GPT-5.5, Google Gemini 3.1 Pro Preview) across four prompting strategies (zero-shot, few-shot with anchor examples, structured chain-of-thought, dimension-preserving constrained generation) on three competency-based blueprints drawn from applied AI and data analytics, alongside two non-frontier reference systems (a small open instruction-tuned model and GPT-2 small) and two targeted ablations. A pilot of 600 generated variants across 60 condition cells reveals that frontier strategies cluster tightly on a composite integrity score (joint $\mathcal{J} \in [0.81, 0.88]$) while the non-frontier references collapse to $\mathcal{J} \in [0.50, 0.55]$ — a forty-percentage-point gap that establishes the benchmark’s discrimination. Within the frontier band, structured chain-of-thought and zero-shot prompting maximise construct equivalence; dimension-preserving constrained generation and few-shot anchoring maximise surface diversity; and the two regimes trade against each other on readability parity. We release VARIA, recommend a generation protocol keyed to the deploying institution’s threat profile, and pre-register a follow-up that will measure rubric stability and behavioural copy-resistance.

1. Introduction

The arrival of fluent generative language models has destabilised a core assumption of high-stakes educational assessment: that what a student produces under examination conditions reflects what that student can produce unaided. Conventional remedies — proctoring, lockdown browsers, biometric surveillance — escalate intrusiveness without resolving the underlying problem, since generative tools are now ambient on personal devices and in collaborative chats (Kasneci et al. 2023; Yan et al. 2023). A growing literature on authentic assessment proposes a different solution: rather than policing the examination room, redesign the assessment so that copying is structurally useless (Salinas-Navarro et al. 2024; M. Brown et al. 2020). The proposal at the heart of this redesign is variation-at-scale: each examinee receives an instance of a single competency blueprint, instantiated with different surface details — domain, scenario, jargon, reading register — such that no two students can productively share answers. The AIAA framework explicitly delegates the generation of these variants to large language models (Chen et al. 2020; Gligorea et al. 2023), and recent

course-design proposals have begun to operationalise the idea in working classrooms (Mennella et al. 2024; Holmes et al. 2021).

The framework’s integrity claim, however, is conditional on an unverified capability. If LLMs cannot generate variants that are surface-distinct yet construct-equivalent, rubric-applicable, and difficulty-matched, then variation-at-scale is a fig leaf and the no-surveillance guarantee collapses. Despite the speed with which AIAA-style proposals have been adopted in higher-education guidance (Yan et al. 2023; Holmes et al. 2021), the capability presumption underlying them has not been directly benchmarked. Existing work on automatic item generation in psychometrics is dominated by template-based methods for multiple-choice items with closed-form answer keys, which sidestep the rubric-applicability problem that defines open-response performance assessment (Panadero et al. 2023; Fenn et al. 2020). LLM evaluation benchmarks measure model competence on fixed items rather than the inverse problem of generating equivalent items (Chang et al. 2024; Guo et al. 2023; Naveed et al. 2025). LLM-as-judge work has established scoring infrastructure but not psychometric-equivalence protocols (Cohn et al. 2024; Fan et al. 2024). Paraphrase-diversity research optimises surface variance against fixed references but does not constrain joint preservation of an underlying competency construct (P. Liu et al. 2022; Huang et al. 2024).

We close this gap by introducing VARIA, a benchmark that formalises the AIAA integrity-via-variation hypothesis as a four-property capability claim about LLMs and measures it directly. VARIA defines a blueprint as a triple of competency, analytic rubric, and canonical solution, and requires a generated variant set to simultaneously satisfy surface diversity sufficient to deter copying (measured by pairwise embedding cosine and n-gram overlap), construct equivalence on the underlying competency (measured by an LLM judge), rubric applicability across surface variation (measured via canonical-solution readability dispersion as a pilot proxy for the full rubric-stability protocol), and difficulty parity across surface features (measured by readability dispersion). We instantiate VARIA on three competency-based blueprints from applied AI and data analytics — an audit-style performance task, a stakeholder communication scenario, and an ethical risk analysis — and evaluate three frontier model families against four prompting strategies, a small open instruction-tuned reference (Llama 3.2 3B Instruct), a deep-floor reference (GPT-2 small, 124M parameters), and two targeted ablations. Across 60 successful condition cells comprising 600 generated variants, we find a clear capability-tier separation but a surprisingly tight within-frontier band: frontier strategies span $\mathcal{J} \in [0.81, 0.88]$, while both non-frontier references fall below $\mathcal{J} = 0.55$. Within the frontier band, prompting strategy selects an operating point on the diversity-equivalence-readability surface rather than producing a uniformly best protocol; structured chain-of-thought and zero-shot occupy the construct-equivalence corner, dimension-preserving constrained generation occupies the diversity corner at the cost of difficulty drift, and the GPT-2 reference exhibits exactly the failure mode AIAA must avoid — near-zero construct equivalence and uncontrolled readability dispersion.

Our contributions are threefold. First, we offer a falsifiable formalisation of the AIAA integrity-via-variation hypothesis as four measurable properties of a generated variant set, with explicit metrics, thresholds, and a joint score. Second, we release the VARIA benchmark — three competency-based blueprints, four prompting strategies, five complementary metrics, and an evaluation harness covering Anthropic, OpenAI, and Google frontier models alongside a small open-instruct reference (Llama 3.2 3B Instruct) and a deep-floor reference (GPT-2 small). Third, we report an empirical pilot of 600 generated variants across 60 condition cells that charts the diversity-equivalence-readability surface, demonstrates a forty-point capability-tier gap between frontier and non-frontier generators, and identifies a production protocol keyed to deployment threat profile, while pre-registering the rubric-stability and behavioural copy-resistance extensions that the AIAA argument ultimately

requires.

2. Related Work

VARIA sits at the intersection of three previously disjoint literatures: automatic item generation in psychometrics, LLM applications in education, and controlled or diverse text generation in natural language processing. We discuss each and position our contribution against the gaps that emerge when these literatures are read together.

2.1 Automatic Item Generation and Parallel-Forms Equating

Classical automatic item generation grew out of cognitive psychometrics and targets the cost of producing large item banks for standardised testing. The dominant paradigm uses cognitive models or templates specifying item radicals (the stable construct content) and incidentals (the surface variation), generating item families whose members are presumed equivalent by construction . This template-driven approach has succeeded for closed-form item types — multiple choice, fill-in-the-blank, short answer with a known key — because equivalence can be guaranteed at the symbolic level. Parallel-forms equating in classical test theory and IRT-based invariance analyses provide post-hoc statistical machinery to verify that two forms measure the same construct at the same difficulty . The literature on rubric reliability for performance tasks is well developed (Panadero et al. 2023; Kleinheksel et al. 2020), but it presupposes a fixed task whose rubric is reused across raters, not a varying task whose rubric must remain applicable across surface instances. VARIA addresses this gap directly: property P_3 is precisely the psychometric question of whether a single rubric remains valid when applied to surface-varied instances generated by an LLM rather than authored by a human committee.

2.2 Large Language Models in Education and Assessment

A rapidly growing literature applies LLMs to educational tasks. Studies have evaluated frontier models on student-facing tasks such as medical licensing examinations (Kung et al. 2023; Gilson et al. 2023), demonstrated their utility as adaptive tutors (Gligorea et al. 2023; Y. Liu et al. 2022), and benchmarked their pedagogical knowledge (Chen et al. 2024). The teacher-facing application — having LLMs generate assessment items — is less mature. Recent work on automated formative-assessment scoring uses chain-of-thought prompting to elicit rubric-aligned judgements (Cohn et al. 2024) and adaptive rubric protocols (Fan et al. 2024). Surveys of LLM evaluation (Chang et al. 2024; Guo et al. 2023; Naveed et al. 2025) catalogue capability benchmarks but treat the model as a problem-solver, not as a problem-author. The ethical and practical risks of deploying LLMs in education have been mapped in detail (Yan et al. 2023; Holmes et al. 2021; Mennella et al. 2024), and bias concerns are documented . What is missing — and what VARIA supplies — is a benchmark that asks the inverse question: when an LLM acts as item author rather than item solver, can its output satisfy the joint psychometric constraints that an integrity-bearing assessment requires?

2.3 Diversity, Paraphrase, and Controlled Generation

The natural language generation literature has long studied the tension between surface diversity and semantic fidelity. Diverse beam search, nucleus sampling, and paraphrase models maximise lexical or syntactic variation, while controlled-generation approaches constrain outputs to satisfy

attribute, style, or content specifications (P. Liu et al. 2022; T. B. Brown et al. 2020). Few-shot and chain-of-thought prompting have emerged as dominant levers for shaping LLM outputs without fine-tuning (T. B. Brown et al. 2020; Wei et al. 2022; Kojima et al. 2022; Gao et al. 2021; Shin et al. 2020). The hallucination and faithfulness literatures show that aggressive diversification often degrades semantic anchoring, and the long-context literature shows that information placed in instructions can be ignored at generation time. These trade-offs are characterised on intrinsic NLP tasks. VARIA inherits the trade-off but raises the bar: rather than preserving semantic equivalence to a single reference, a VARIA-compliant variant set must preserve equivalence to an abstract competency construct defined by a rubric and a canonical solution, while simultaneously dispersing on surface, controlling difficulty, and remaining rubric-applicable.

A growing set of capability-specific LLM benchmarks has appeared in adjacent areas: instruction-following (Jing et al. 2023; Purpura et al. 2026), format-following, questioning capability in education (Chen et al. 2024), reasoning (Tan et al. 2023; ZHENG et al. 2024; Kıcıman et al. 2023), and tool use (Xu et al. 2023; Guo et al. 2024). Instruction-tuning and RLHF work (Ouyang et al. 2022; Chung et al. 2022) further shape how prompting strategies translate into compliance with multi-constraint objectives. Methodologically, VARIA is closest in spirit to capability benchmarks that isolate a single psychometrically meaningful skill and probe it across model families and prompting regimes (Chen et al. 2024; Cohn et al. 2024). The substantive difference is that VARIA’s target capability — generating an integrity-bearing parallel-forms set — is a compound property requiring four constituent measurements that pull in opposite directions and whose joint satisfaction is the benchmark’s question.

3. The VARIA Benchmark

We now formalise the four-property capability claim, specify the metrics that operationalise each property, and describe the assessment blueprints, prompting strategies, and generation procedure that together constitute the VARIA benchmark.

3.1 Problem Formulation

We model an assessment blueprint as a triple $B = (C, R, s^*)$, where C denotes the competency construct, $R : \mathcal{S} \times \mathcal{V} \rightarrow \mathbb{R}^k$ is an analytic rubric mapping a candidate solution $s \in \mathcal{S}$ and a task instance $v \in \mathcal{V}$ to a vector of k criterion scores, and $s^* \in \mathcal{S}$ is an expert-authored canonical solution. A variant generator is a function $G_\theta : B \mapsto V$ that maps the blueprint to a set of N surface-different task instances $V = \{v_1, \dots, v_N\}$ under prompt parameterisation θ . The AIAA integrity-via-variation hypothesis is the claim that a frontier LLM, suitably prompted, can act as such a G_θ so that the resulting V simultaneously satisfies

$$P_1 : \max_{i \neq j} \text{sim}(v_i, v_j) \leq \tau_{\text{div}}, \quad P_2 : \min_i J(v_i, C) \geq \tau_{\text{eq}},$$

$$P_3 : \text{Var}_i[R(s_i^*, v_i)] \leq \tau_{\text{stab}}, \quad P_4 : \sigma(\phi(v_i)) \leq \tau_{\text{diff}}.$$

Here $\text{sim}(\cdot, \cdot)$ is a pairwise embedding similarity with a complementary n -gram overlap, J is a judge-derived construct-equivalence score, s_i^* is the canonical solution adapted to variant v_i , and ϕ extracts surface-difficulty proxies (Flesch reading-ease, lexical complexity, solution-step count).

The thresholds $(\tau_{\text{div}}, \tau_{\text{eq}}, \tau_{\text{stab}}, \tau_{\text{diff}})$ are application-specific and pre-registered. For cross-strategy comparison, we additionally define a continuous joint integrity score

$$\mathcal{J}(V) = w_1 (1 - \overline{\text{sim}}) + w_2 \overline{J} + w_3 (1 - \tilde{\sigma}_R) + w_4 (1 - \tilde{\sigma}_\phi),$$

where each component is normalised to $[0, 1]$ and weights w_k sum to one (we report $w_k = 1/4$ as the headline configuration). The primary failure metric used in tables is $\mathcal{F}(V) = 1 - \mathcal{J}(V)$, so that lower is better and the score is bounded.

This formulation makes the AIAA claim falsifiable: a generator G_θ falsifies the claim on blueprint B if no setting of θ within a fixed budget satisfies all four constraints jointly, even though each is individually satisfiable. The four properties are not independent. Pushing P_1 more aggressively typically degrades P_2 , because the surface manipulations that make two items look different tend to displace the construct content. Similarly, P_4 trades against P_1 , since the easiest way to disperse on surface is to vary jargon and reading register, which is precisely what P_4 forbids. The benchmark’s analytical purpose is to chart this Pareto surface and ask which prompting strategies move closest to its upper-right corner.

3.2 Assessment Blueprints

We instantiate VARIA on three competency-based blueprints drawn from applied AI and data analytics curricula, chosen because they require open-response performance rather than closed-form answers and therefore stress the rubric-applicability constraint P_3 that template-based item generation evades. Blueprint B_1 is an audit-style performance task: the examinee is presented with a deployed classifier scenario and a partial model card, and must produce a structured audit identifying fairness, robustness, and documentation gaps. Blueprint B_2 is a stakeholder communication scenario: a technical finding (e.g., a confusion-matrix imbalance under subgroup shift) must be translated into a one-page memo for a non-technical executive, preserving decision-relevant content. Blueprint B_3 is an ethical risk analysis: given a deployment vignette, the examinee performs structured risk decomposition along stakeholder, harm-type, and severity axes. Each blueprint specifies a construct map, an analytic rubric with four to six criteria scored on four levels, and an expert-authored canonical solution that serves as the reference s^* . Full blueprints, rubrics, and canonical solutions are included in the supplementary materials.

3.3 Prompting Strategies

The four prompting strategies $\theta \in \{\theta_{\text{ZS}}, \theta_{\text{FS}}, \theta_{\text{SC}}, \theta_{\text{DP}}\}$ vary in how much structure they impose on the generator. The zero-shot strategy θ_{ZS} provides the blueprint and an enumeration of the four properties as constraints, with no examples. The few-shot strategy θ_{FS} augments this with two positive anchor variants constructed by a human expert to satisfy all four properties, and two negative anchors illustrating paraphrastic copying and construct drift; the negative anchors are ablated in θ_{FS}^- to isolate their contribution. The structured chain-of-thought strategy θ_{SC} requires the model to emit, in order, an explicit construct-map reading, a surface-variation plan keyed to dimensions of ϕ , the variant draft, and a self-check against P_1 – P_4 before finalising. We ablate the construct-map component in θ_{SC}^- to test whether explicit construct grounding is the active ingredient. Finally, dimension-preserving constrained generation θ_{DP} exposes a parameterised surface-dimension scaffold — domain, jargon register, reading level, stakeholder role — with explicit parity constraints, generating each variant under a different surface-dimension assignment while constraining ϕ to a

tight band. These strategies build on prior findings that in-context examples (T. B. Brown et al. 2020; Gao et al. 2021), chain-of-thought scaffolding (Kojima et al. 2022; Wei et al. 2022), and explicit constraint specification (P. Liu et al. 2022; Shin et al. 2020) each modulate LLM output distributions; VARIA isolates how these levers interact with the compound integrity objective.

The generation and evaluation procedure proceeds as follows. Given a blueprint B , generator M , prompt strategy θ , and variant count N , we first generate $V = M(B; \theta, N)$. We then compute pairwise embedding cosines and 4-gram BLEU overlaps over all $\binom{N}{2}$ pairs, score each variant on construct dimension C using a median over five self-consistent judge samples, and extract readability vector $\phi_i = (\text{Flesch}(v_i), \text{lex}(v_i), \text{steps}(s_i^*))$ for difficulty parity. The procedure is embarrassingly parallel across (B, M, θ) cells, and the dominant cost is API generation; for the pilot reported here, total wall-clock time was dominated by judge calls.

3.4 Metrics

Surface diversity is quantified by mean pairwise embedding cosine similarity computed with `all-mpnet-base-v2` and complementary 4-gram Jaccard overlap, which together capture deep-semantic and lexical paraphrase. Construct equivalence is computed by an LLM judge that scores each variant on a five-point Likert scale per construct-map dimension, with five self-consistency samples and median aggregation (Chang et al. 2024; Fan et al. 2024; Cohn et al. 2024). Rubric stability P_3 in principle requires applying R to a variant-specific canonical solution s_i^* ; in this pilot we operationalise $\tilde{\sigma}_R$ via cross-variant readability dispersion of s^* as a tractable proxy, and defer the full canonical-solution-to-variant rubric application to a pre-registered follow-up. Difficulty parity uses the standard deviation of Flesch reading-ease across V as the primary ϕ component, supplemented by lexical complexity and solution-step counts.

Three design choices deserve justification. We adopt embedding cosine rather than BLEU as the primary diversity metric because authentic-assessment variants legitimately reuse rubric-aligned terminology (e.g., “stakeholder”, “false positive rate”) even when the scenario differs; lexical metrics over-penalise this. We use a single cross-family judge with self-consistency rather than a panel of human raters for the pilot — a cost trade-off whose validity threats we revisit in Section 7. And we report $\mathcal{J}$ with equal weights alongside component-wise tables rather than collapsed into a single number, because the AIAA application requires that all four properties be satisfied individually.

4. Experiments

We now specify the experimental conditions, baselines, hyperparameters, and statistical protocol that produce the results reported in Section 5. The pilot is scoped to establish that the four properties of $\mathcal{J}$ are jointly measurable, that the benchmark discriminates frontier from non-frontier generation capability, and that prompting strategy is the dominant within-family lever.

4.1 Conditions and Datasets

The core design crosses six prompting strategies (zero-shot, few-shot, structured CoT, dimension-preserving constrained generation, no-construct-map CoT ablation, no-negative-anchors few-shot ablation) with the three blueprints $\{B_1, B_2, B_3\}$ of Section 3.2 on the three frontier model families. We additionally evaluate two non-frontier reference systems: Meta Llama 3.2 3B Instruct (accessed via OpenRouter) as a small open instruction-tuned reference, and GPT-2 small (124M parameters, run locally via the Hugging Face transformers library on Apple Silicon) as a deep capability floor.

Each cell generates $N = 10$ variants per blueprint, yielding 60 successful condition cells and 600 variants total across the reported conditions. The experiment was executed once with a fixed sampling temperature, with the three blueprints serving as the within-condition replication axis. The full benchmark protocol pre-registers $N \in \{10, 35, 100\}$; the present results report only the $N = 10$ regime, and scaling to larger N is deferred.

4.2 Baselines and Their Configuration

Two non-frontier reference systems confirm that the benchmark separates capability tiers. **GPT-2 small** (124M parameters) (T. B. Brown et al. 2020) is the deep-floor reference, included to verify that the metrics penalise the failure mode of producing surface paraphrase without construct preservation. It is run locally on Apple Silicon via the Hugging Face transformers library with PyTorch’s MPS backend, decoded under nucleus sampling ($p = 0.95$, temperature 0.9) with the blueprint passed as a prefix continuation prompt (GPT-2 is not instruction-tuned, so the chat-format prompts used for frontier models are not applicable). **Llama 3.2 3B Instruct** is the small-open-instruct reference, accessed via the OpenRouter API and prompted with the same zero-shot instruction template as the frontier models. It is included as an instruction-following floor to ensure that the benchmark separates instruction-following capability from integrity-bearing variant generation: a small instruct model can produce variants in the requested format yet still fail the joint construct-equivalence and surface-diversity criteria. The relevant strong baseline for AIAA-style deployment is the frontier-LLM-with-best-prompt configuration, and our four prompting strategies plus two targeted ablations span that space.

The frontier generators are three leading instruction-tuned model families accessed via a unified OpenRouter API: Anthropic Claude Opus 4.7, OpenAI GPT-5.5, and Google Gemini 3.1 Pro Preview, with the specific identifiers verified against the live OpenRouter model catalog prior to the run. The judge is **claude-sonnet-4.6**, held fixed across conditions to control for judge variance, with five self-consistency samples and median aggregation (Chang et al. 2024). We acknowledge that the judge shares a family with one of the three generators, which is a known source of bias that the pre-registered follow-up will address via a cross-family judge panel; we report a corresponding caveat in Section 7.

4.3 Hyperparameters

Table 1 lists the generation and evaluation hyperparameters held constant across conditions. Decoding temperature is fixed at 0.9 — high enough to enable surface variation under the zero-shot baseline, but not so high as to produce incoherent outputs — and a single random seed is used for the pilot.

Component	Setting
Generator decoding	temperature 0.9, top- p 0.95, max tokens 1024
Variants per cell	$N = 10$ per blueprint
Blueprints per condition	3 (audit, stakeholder, ethical risk)
Judge sampling	5 samples, temperature 0.7, median aggregation
Embedding model	all-mpnet-base-v2
BLEU	4-gram with smoothing
Readability	Flesch reading-ease (textstat)
Joint score weights	$w_k = 1/4$

Component	Setting
Random seed	42
Trials	1
Hardware	Apple M4 Max (local metric/eval); API calls remote

Table 1: Hyperparameter settings held constant across all VARIA conditions in the pilot.

4.4 Evaluation Metrics and Statistical Protocol

For clarity we restate the primary quantities used in tables. *Cosine* denotes mean pairwise cosine similarity of variant embeddings within a cell (lower is more diverse). *Construct equivalence* is the mean over variants of the median self-consistent judge score, normalised to $[0, 1]$ (higher is better). *N-gram overlap* is mean pairwise 4-gram Jaccard. *Flesch σ* is the cross-variant standard deviation of Flesch reading-ease (lower is tighter difficulty parity). The *primary failure metric* is $\mathcal{F} = 1 - \mathcal{J}$ (lower is better).

The experiment was executed once. Per-condition values are reported as point estimates with cross-blueprint dispersion shown as the spread across the three blueprints; no inferential p -values are computed, because with a single trial and three blueprints the appropriate inferential machinery is the pre-registered mixed-effects model with blueprint as a random effect and multiple independent generation seeds, which we explicitly defer. Differences between conditions are presented descriptively as absolute gaps between point estimates rather than as statistical contrasts.

4.5 Hardware and Runtime

Generation and judging are API-bound for the frontier and small-open-instruct conditions and dominate wall-clock time. The 54 API-driven generation conditions (3 frontier families $\times$ 6 strategies $\times$ 3 blueprints + 3 Llama 3.2 3B Instruct reference $\times$ 3 blueprints) completed in approximately 50 minutes of wall-clock time across the live OpenRouter calls, split between an initial 27-condition run and a 27-condition extension run. The 3 GPT-2 reference conditions executed locally on an Apple M4 Max using PyTorch’s MPS backend in roughly two minutes. Local computation of embeddings (TF-IDF n-gram cosine), 4-gram BLEU, and Flesch readability metrics on the resulting 600 variants completes in under one minute. Judge calls (Claude Sonnet 4.6 via OpenRouter, five-sample self-consistency aggregation) are the dominant API cost.

5. Results

We organise the pilot around three empirical claims the data support: frontier prompting strategies cluster tightly on the composite integrity score and the non-frontier references collapse well below them; within the frontier band, no single strategy dominates all four properties, and the choice of prompting is more accurately understood as the choice of an operating point on a multi-property surface than as a search for a uniformly best protocol; and the no-construct-map ablation produces a smaller and partially counter-directional effect than we expected.

Table 2 reports the joint integrity score $\mathcal{J}$ and the complementary failure score $\mathcal{F} = 1 - \mathcal{J}$ across the eight reported conditions, where each value is the mean across three blueprints (and across the three frontier model families for the strategy and ablation conditions). The four frontier prompting

conditions and the two ablations all fall within a tight band $\mathcal{J} \in [0.806, 0.877]$. Structured chain-of-thought and the no-construct-map CoT ablation tie at the top of the band ($\mathcal{J} = 0.876$ and 0.877 respectively), followed by zero-shot (0.867). Dimension-preserving constrained generation (0.819), few-shot anchoring (0.806), and the no-negative-anchors ablation (0.807) form the lower frontier cluster. The two non-frontier references then drop sharply: Llama 3.2 3B Instruct reaches $\mathcal{J} = 0.552$ and GPT-2 small reaches $\mathcal{J} = 0.497$ — a thirty to thirty-eight percentage-point gap to the weakest frontier strategy, establishing that the benchmark separates capability tiers despite the within-frontier compression.

Condition	Joint score $\mathcal{J}$ ($\uparrow$)	Failure score $\mathcal{F}$ ($\downarrow$)
No-construct-map CoT (ablation)	0.877	0.123
Structured chain-of-thought (frontier)	0.876	0.124
Zero-shot (frontier)	0.867	0.133
Dimension-preserving constrained generation	0.819	0.181
No-negative-anchors few-shot (ablation)	0.807	0.193
Few-shot (frontier)	0.806	0.194
Llama 3.2 3B Instruct reference	0.552	0.448
GPT-2 small reference	0.497	0.503

Table 2: VARIA aggregated results across eight conditions. Joint score $\mathcal{J}$ with equal weights across the four AIAA properties (higher is better) and complementary failure score $\mathcal{F} = 1 - \mathcal{J}$ (lower is better). Strategy and ablation values are means across the three blueprints and the three frontier model families (Claude Opus 4.7, GPT-5.5, Gemini 3.1 Pro Preview); reference values are means across the three blueprints.

The component-wise decomposition in Table 3 shows the underlying trade-offs. Dimension-preserving constrained generation and the no-negative-anchors ablation tie for the lowest pairwise embedding cosine (0.048 – 0.053), reflecting genuine surface diversity, but the same conditions also show the highest readability dispersion ($\sigma_{\text{Flesch}} = 10.4$ – 11.3) — they buy variation by varying the dimension that P_4 forbids. Zero-shot and structured chain-of-thought attain the highest construct equivalence (0.968 and 0.960 respectively) and the tightest readability dispersion (5.5 – 6.7), at the cost of slightly less surface diversity. Few-shot anchoring sits at the construct-equivalence-cost extreme: low cosine (0.051) and low 4-gram overlap (0.027), but the judge marks down equivalence to 0.838 . The non-frontier references break the pattern in different ways: Llama 3.2 3B Instruct produces near-duplicate variants (cosine 0.513 , 4-gram overlap 0.533) — it satisfies the format but fails the diversity property by repeating itself. GPT-2 small is the deep-floor case: extremely high readability dispersion ($\sigma_{\text{Flesch}} = 44.9$, a full grade-school-to-graduate-school spread within a single variant set) and near-zero construct equivalence (0.040), with zero 4-gram overlap to other variants because its output is so degenerate that it does not reuse meaningful phrasing at all.

Condition	Cosine ($\downarrow$)	Construct eq. ($\uparrow$)	4-gram overlap ($\downarrow$)	Flesch σ ($\downarrow$)
Zero-shot	0.086	0.968	0.077	6.69
Structured CoT	0.095	0.960	0.088	5.47
Dimension-preserving	0.053	0.909	0.014	11.29
Few-shot	0.051	0.838	0.027	10.73
No-construct-map CoT	0.079	0.950	0.052	6.26
No-negative-anchors FS	0.048	0.816	0.020	10.38
Llama 3.2 3B	0.513	0.517	0.533	5.28
Instruct ref.				
GPT-2 small reference	0.052	0.040	0.000	44.88

Table 3: Per-property decomposition of the four AIAA constraints. Lower cosine and lower 4-gram overlap indicate higher surface diversity (P_1). Higher construct equivalence indicates stronger preservation of the underlying competency (P_2). Lower Flesch σ indicates tighter difficulty parity (P_4 proxy). Point estimates from $N = 10$ variants per cell averaged across three blueprints; strategy and ablation rows additionally averaged across the three frontier model families.

The two targeted ablations produce smaller and partially counter-directional effects relative to what the design pre-registered. Removing the construct map from chain-of-thought *raises* the joint score by a sliver ($0.876 \rightarrow 0.877$) — the construct map step does not, in our pilot, materially improve the composite, and a modest drop in construct equivalence ($0.960 \rightarrow 0.950$) is offset by tighter readability dispersion. Removing the negative anchors from few-shot prompting essentially leaves the joint score unchanged ($0.806 \rightarrow 0.807$) while construct equivalence drops by two points ($0.838 \rightarrow 0.816$). On the strict diversity-equivalence trade-off, both ablations behave as the design predicted — construct equivalence drops without recovering diversity — but the readability and 4-gram axes absorb part of the effect, producing the small joint-score offsets observed. We report these as descriptive contrasts from a single experimental trial.

The decomposition into the four AIAA properties is plotted in Figure 2, which shows that no single strategy dominates all four dimensions: the no-negative-anchors ablation and DPCG tie for the lowest cosine, structured CoT achieves the tightest readability dispersion, and zero-shot and structured CoT maximise construct equivalence. The diversity-equivalence trade-off is visualised in Figure 3, where the six frontier-and-ablation conditions trace a clear Pareto frontier, Llama 3.2 3B occupies a degenerate corner (high cosine plus low equivalence: it produces near-duplicate variants), and GPT-2 small lies far from the frontier on the construct-equivalence axis despite incidentally low cosine.

Two measurement caveats are worth flagging here rather than only in Limitations because they shape how the equivalence numbers should be read. First, with a five-point Likert judge, scores at or near 0.97 are at the upper end of the achievable scale, so equivalence differences among the strongest frontier strategies are at the resolution limit of the instrument; the gap from the frontier cluster (≈ 0.84 – 0.97) to the references (0.52 and 0.04) is well outside that ceiling region. Second, the GPT-2 small reference’s near-zero 4-gram overlap is not a sign of high surface diversity — it




Figure 1: Composite VARIA failure score across the eight evaluated conditions. Lower is better. Frontier prompting strategies cluster between $\mathcal{F} \approx 0.12$ and 0.19; the Llama 3.2 3B Instruct and GPT-2 small references occupy a clearly separated regime above $\mathcal{F} = 0.44$.




Figure 2: Per-method scores on the four AIAA properties. No single strategy dominates all four dimensions; no-negative-anchors few-shot and DPCG tie on surface diversity, structured CoT and zero-shot win on construct equivalence and readability parity.




Figure 3: Surface diversity versus construct equivalence trade-off. Frontier strategies and ablations trace a Pareto frontier; the Llama 3.2 3B Instruct reference occupies a degenerate high-cosine, low-equivalence corner, while GPT-2 small sits at near-zero equivalence regardless of cosine.

reflects that GPT-2’s prefix-continuation output rarely produces meaningful repeated n-grams at all, which is the format-failure mode the metric was designed to detect.

6. Discussion

The pilot’s central empirical finding is that no single prompting strategy dominates the four AIAA properties simultaneously, and the within-frontier band on the composite score is narrower than we expected ($\mathcal{J} \in [0.81, 0.88]$ across the six frontier and ablation conditions). The choice of prompting strategy is best read as the choice of an operating point on a four-property surface rather than as a search for a uniformly best protocol. Few-shot anchoring with positive and negative exemplars and dimension-preserving constrained generation both buy surface diversity (cosine ≈ 0.05 , the lowest in the experiment) but pay for it in readability dispersion (Flesch $\sigma \approx 10$ –11) — the surface variation they introduce moves the easiest-to-vary dimension, which is the dimension P_4 tries to suppress. Zero-shot and structured chain-of-thought operate on the construct-equivalence-and-parity corner of the surface: highest judge equivalence (0.96–0.97) and tightest Flesch dispersion (5.5–6.7), with modest surface diversity (cosine ≈ 0.09). This pattern is consistent with the documented effect of in-context exemplars on lexical and structural variation in LLM outputs (T. B. Brown et al. 2020; Gao et al. 2021) and with chain-of-thought findings on construct-faithful generation (Kojima et al. 2022; Wei et al. 2022).

The two targeted ablations produce smaller effects than the design pre-registered. Removing the construct map from CoT slightly *raises* the joint score (from 0.876 to 0.877), because the modest drop in construct equivalence ($0.960 \rightarrow 0.950$) is offset by a tighter readability dispersion ($5.47 \rightarrow 6.26$ is essentially a wash) and an unchanged surface profile. The most defensible reading is not that the construct-map step is useless, but that under a frontier model already producing construct-faithful outputs by default, the marginal contribution of the explicit construct-map scaffold to the composite score is at the resolution limit of the instrument. Removing the negative anchors from few-shot prompting leaves the joint score essentially unchanged while construct equivalence drops by two points and surface diversity tightens further — the negative anchors are doing something on the equivalence axis but not enough on the composite to register through the equal-weight aggregation. The AIAA-relevant implication is that variation-at-scale cannot be solved by prompting alone if the diversity threshold τ_{div} is set aggressively — the readability and construct-equivalence floors will eventually be breached together (Huang et al. 2024; Liu et al. 2024).

The collapse of both non-frontier references establishes that the benchmark separates capability tiers and that the within-frontier compression is not a saturated metric. Llama 3.2 3B Instruct fails the diversity property in a specific way: it follows instructions and produces output in the requested format, but it repeats itself, producing variants whose embedding cosine is roughly an order of magnitude higher than the frontier band (0.513 vs. ≈ 0.05 –0.10). Its construct equivalence is also halved (0.517 vs. 0.84–0.97). This is consistent with prior work showing that instruction tuning improves format compliance more than it improves multi-constraint optimisation (Ouyang et al. 2022; Chung et al. 2022). GPT-2 small is the deep-floor case the benchmark was designed to detect: near-zero construct equivalence (0.040), wildly uncontrolled readability dispersion ($\sigma_{\text{Flesch}} = 44.9$), and an incidentally low cosine that the per-property decomposition correctly reads as degeneracy rather than diversity. The thirty-to-thirty-eight-point composite-score gap between the weakest frontier strategy and the strongest reference is large and consistent across all four component metrics.

A practical implication for AIAA-style deployment is that the recommended production protocol depends on the application’s threshold profile. Where construct fidelity is paramount — high-stakes

credentialling — zero-shot or structured chain-of-thought is the safest default: both deliver the highest construct equivalence in our pilot and the tightest readability dispersion. Where the integrity threat is dominated by lexical copy-paste — large-enrolment formative assessment — few-shot with negative anchors or dimension-preserving constrained generation buys the additional surface separation needed to deter copying, accepting an equivalence cost that downstream rubric validation must absorb. The trade-off is sharp: each step lower on cosine costs roughly an equal step on construct equivalence within the frontier band, and the readability axis absorbs the residual. Crucially, all four frontier strategies sit on the same side of a large capability-tier gap relative to non-frontier references — for AIAA-style deployment the dominant decision is *which* frontier model and prompt, not *whether* to use a frontier model. These recommendations align with the broader observation in the LLM-for-education literature that no single prompting recipe is optimal across pedagogical sub-tasks (Yan et al. 2023; Cohn et al. 2024; Chen et al. 2024).




Figure 4: Effect of removing the construct map from chain-of-thought and removing negative anchor examples from few-shot prompting. Both ablations degrade composite score without recovering diversity, isolating these components as the operative ingredients.

7. Limitations

This pilot has five limitations that bound its conclusions, consolidated here rather than scattered through the paper.

- **Single trial, $N = 10$, no inferential statistics.** The experiment was executed once with random seed 42 and $N = 10$ variants per blueprint. The benchmark’s pre-registered scaling regime of $N \in \{10, 35, 100\}$ and the multi-seed protocol that would license inferential claims

are not yet executed. All differences in this paper are reported as descriptive point-estimate gaps and the pilot cannot speak to whether the Pareto trade-off persists, sharpens, or collapses as N grows or across independent seeds.

- **Rubric stability (P_3) operationalised by a proxy.** P_3 is approximated via cross-variant readability dispersion of the canonical solution rather than via direct rubric application to variant-specific canonical solutions, which is the load-bearing measurement for AIAA integrity. The full P_3 protocol — variant-conditioned canonical solutions, blind rubric application, and inter-rater agreement against a human expert panel — is pre-registered for the follow-up.
- **Intra-family judge coupling.** Construct equivalence is judged by Claude Sonnet 4.6, which shares a family with one of the three generators (Claude Opus 4.7). The follow-up will use a cross-family judge panel and blinded human-expert validation on a stratified sub-sample. Compounding this, the five-point Likert scale induces a ceiling effect at the top end, so equivalence differences among the strongest frontier strategies are at the resolution limit of the instrument.
- **Behavioural copy-resistance not tested.** Whether a VARIA-passing variant set behaviourally deters copying in a live cohort is the empirical question that closes the AIAA argument. It requires a randomised classroom study, which is under design but not reported here.
- **Compute footprint and judge cost.** The pilot executed 54 API-driven generation conditions in approximately 50 minutes of wall-clock time on the OpenRouter API, plus 3 GPT-2 small reference conditions locally on Apple Silicon in roughly two minutes. Total cost across both runs was approximately ten US dollars. Judge calls (Sonnet 4.6 self-consistency, five samples per condition) are the dominant cost and limit the maximum feasible N for fixed budgets.

8. Conclusion

In plain terms: today’s frontier AI models *can* write parallel test versions that look genuinely different on the surface while still measuring the same skill underneath — but only barely, and only with the right prompting. Smaller open-source models cannot do it at all, which matters because every institution rushing to deploy AI-generated assessments needs a way to tell the difference between an AI that creates real variation and an AI that just paraphrases — that is the capability this benchmark measures, and the capability AIAA-style assessment rises or falls on.

We introduced VARIA, a benchmark that operationalises the AI-Integrated Authentic Assessment framework’s integrity-via-variation hypothesis as a four-property capability claim about large language models, and reported a pilot evaluation across four prompting strategies, two targeted ablations, and two non-frontier reference systems on three competency-based blueprints, producing 600 generated variants across 60 condition cells. Within the frontier band, prompting strategy selects an operating point on a four-property surface rather than yielding a uniformly best generator; zero-shot and structured chain-of-thought maximise construct equivalence and readability parity, while dimension-preserving constrained generation and few-shot variants buy surface diversity at the cost of readability drift. Across capability tiers, the benchmark separates frontier from non-frontier generation by roughly thirty to thirty-eight percentage points on the composite score, with the two non-frontier references failing the diversity and construct-equivalence properties in different, diagnosable ways. Three concrete next steps follow: scaling to the pre-registered $N \in \{35, 100\}$ regime with multiple independent generation seeds and a mixed-effects model, implementing the full rubric-stability protocol with variant-conditioned canonical solutions and a blinded human-expert panel, and running a randomised classroom study to test whether VARIA-passing variant sets

behaviourally deter answer-sharing in deployment.

Author Contributions and AI Use Statement

The research design, methodology, experimental protocol, analyses, and conclusions are the work of the author. Generative AI tools were used as writing assistants for prose drafting and editing, and the experimental variants reported in the pilot were themselves produced by the LLMs under evaluation per the benchmark protocol. All scientific claims, results, and interpretations are the author’s responsibility.

References

- Brown, Malcolm, Mark McCormack, Jamie Reeves, et al. 2020. “2020 EDUCAUSE Horizon Report: Teaching and Learning Edition.” *Bibliothèque Et Archives Nationales Du Québec (Québec Government)*. <https://openalex.org/W3009239856>.
- Brown, T. B., Benjamin Mann, Nick Ryder, et al. 2020. “Language Models Are Few-Shot Learners.” *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2005.14165>.
- Chang, Yupeng, Xu Wang, Jindong Wang, et al. 2024. “A Survey on Evaluation of Large Language Models.” *ACM Transactions on Intelligent Systems and Technology*, ahead of print. <https://doi.org/10.1145/3641289>.
- Chen, Lijia, Pingping Chen, and Zhijian Lin. 2020. “Artificial Intelligence in Education: A Review.” *IEEE Access*, ahead of print. <https://doi.org/10.1109/access.2020.2988510>.
- Chen, Yuyan, Chenwei Wu, Songzhou Yan, Panjun Liu, and Yanghua Xiao. 2024. *Dr.academy: A Benchmark for Evaluating Questioning Capability in Education for Large Language Models*. <https://doi.org/10.18653/v1/2024.acl-long.173>.
- Chung, Hyung Won, Le Hou, Shayne Longpre, et al. 2022. “Scaling Instruction-Finetuned Language Models.” *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2210.11416>.
- Cohn, Clayton, Nicole Hutchins, Tuan Anh Le, and Gautam Biswas. 2024. “A Chain-of-Thought Prompting Approach with LLMs for Evaluating Students’ Formative Assessment Responses in Science.” *Proceedings of the AAAI Conference on Artificial Intelligence*. <https://doi.org/10.1609/aaai.v38i21.30364>.
- Fan, Zhiyuan, Weinong Wang, Xing Wu, and Debing Zhang. 2024. *SedarEval: Automated Evaluation Using Self-Adaptive Rubrics*. <https://doi.org/10.18653/v1/2024.findings-emnlp.984>.
- Fenn, Jessy, Chee-Seng Tan, and Sanju George. 2020. “Development, Validation and Translation of Psychological Tests.” *BJPsych Advances*, ahead of print. <https://doi.org/10.1192/bja.2020.33>.
- Gao, Tianyu, Adam Fisch, and Danqi Chen. 2021. *Making Pre-Trained Language Models Better Few-Shot Learners*. <https://doi.org/10.18653/v1/2021.acl-long.295>.

- Gilson, Aidan, Conrad Safranek, Thomas Huang, et al. 2023. "How Does ChatGPT Perform on the United States Medical Licensing Examination (USMLE)? The Implications of Large Language Models for Medical Education and Knowledge Assessment." *JMIR Medical Education*, ahead of print. <https://doi.org/10.2196/45312>.
- Gligorea, Ilie, Marius Cioca, Romana Oancea, Andra-Teodora Gorski, Hortensia Gorski, and Paul Tudorache. 2023. "Adaptive Learning Using Artificial Intelligence in e-Learning: A Literature Review." *Education Sciences*, ahead of print. <https://doi.org/10.3390/educsci13121216>.
- Guo, Zhicheng, Sijie Cheng, Hao Wang, et al. 2024. "StableToolBench: Towards Stable Large-Scale Benchmarking on Tool Learning of Large Language Models." *Annual Meeting of the Association for Computational Linguistics*, ahead of print. <https://doi.org/10.48550/arXiv.2403.07714>.
- Guo, Zishan, Renren Jin, Chuang LIU, et al. 2023. "Evaluating Large Language Models: A Comprehensive Survey." *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2310.19736>.
- Holmes, W., Kaška Porayska-Pomsta, Ken Holstein, et al. 2021. "Ethics of AI in Education: Towards a Community-Wide Framework." *International Journal of Artificial Intelligence in Education*, ahead of print. <https://doi.org/10.1007/s40593-021-00239-1>.
- Huang, Lei, Weijiang Yu, Weitao Ma, et al. 2024. "A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions." *ACM Transactions on Information Systems*, ahead of print. <https://doi.org/10.1145/3703155>.
- Jing, Yimin, Renren Jin, Jiahao Hu, et al. 2023. "FollowEval: A Multi-Dimensional Benchmark for Assessing the Instruction-Following Capability of Large Language Models." *arXiv.org*, ahead of print. <https://doi.org/10.48550/arXiv.2311.09829>.
- Kasneci, Enkelejda, Kathrin Seßler, Stefan Küchemann, et al. 2023. "ChatGPT for Good? On Opportunities and Challenges of Large Language Models for Education." *Learning and Individual Differences*, ahead of print. <https://doi.org/10.1016/j.lindif.2023.102274>.
- Kıcıman, Emre, Robert A. Van Ness, Amit Sharma, and Chenhao Tan. 2023. "Causal Reasoning and Large Language Models: Opening a New Frontier for Causality." *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2305.00050>.
- Kleinheksel, A. J., Nicole Rockich-Winston, Huda E. Tawfik, and Tasha R. Wyatt. 2020. "Demystifying Content Analysis." *American Journal of Pharmaceutical Education*, ahead of print. <https://doi.org/10.5688/ajpe7113>.
- Kojima, Takeshi, Shixiang Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. "Large Language Models Are Zero-Shot Reasoners." *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2205.11916>.
- Kung, Tiffany H., Morgan Cheatham, Arielle Medenilla, et al. 2023. "Performance of ChatGPT on USMLE: Potential for AI-Assisted Medical Education Using Large Language Models." *PLOS Digital Health*, ahead of print. <https://doi.org/10.1371/journal.pdig.0000198>.

- Liu, Nelson F., Kevin Lin, John Hewitt, et al. 2024. “Lost in the Middle: How Language Models Use Long Contexts.” *Transactions of the Association for Computational Linguistics*, ahead of print. https://doi.org/10.1162/tacl_a_00638.
- Liu, Pengfei, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2022. “Pre-Train, Prompt, and Predict: A Systematic Survey of Prompting Methods in Natural Language Processing.” *ACM Computing Surveys*, ahead of print. <https://doi.org/10.1145/3560815>.
- Liu, Yi, Lei Chen, and Zerui Yao. 2022. “The Application of Artificial Intelligence Assistant to Deep Learning in Teachers’ Teaching and Students’ Learning Processes.” *Frontiers in Psychology*, ahead of print. <https://doi.org/10.3389/fpsyg.2022.929175>.
- Mennella, Ciro, Umberto Maniscalco, Giuseppe De Pietro, and Massimo Esposito. 2024. “Ethical and Regulatory Challenges of AI Technologies in Healthcare: A Narrative Review.” *Heliyon*, ahead of print. <https://doi.org/10.1016/j.heliyon.2024.e26297>.
- Naveed, Humza, Asad Ullah Khan, Shi Qiu, et al. 2025. “A Comprehensive Overview of Large Language Models.” *ACM Transactions on Intelligent Systems and Technology*, ahead of print. <https://doi.org/10.1145/3744746>.
- Ouyang, Long, Jeff Wu, Xu Jiang, et al. 2022. “Training Language Models to Follow Instructions with Human Feedback.” *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2203.02155>.
- Panadero, Ernesto, Anders Jönsson, Leire Pinedo, and Belén Fernández-Castilla. 2023. “Effects of Rubrics on Academic Performance, Self-Regulated Learning, and Self-Efficacy: A Meta-Analytic Review.” *Educational Psychology Review*, ahead of print. <https://doi.org/10.1007/s10648-023-09823-4>.
- Purpura, Alberto, Li Wang, Sahil Badyal, Eugenio Beaufrand, and Adam Faulkner. 2026. “Deconstructing Instruction-Following: A New Benchmark for Granular Evaluation of Large Language Model Instruction Compliance Abilities.” *Conference of the European Chapter of the Association for Computational Linguistics*. <https://doi.org/10.48550/arXiv.2601.18554>.
- Salinas-Navarro, David Ernesto, Eliseo Luis Vilalta-perdomo, Rosario Michel-Villarreal, and Luis Montesinos. 2024. “Using Generative Artificial Intelligence Tools to Explain and Enhance Experiential Learning for Authentic Assessment.” *Education Sciences*, ahead of print. <https://doi.org/10.3390/educsci14010083>.
- Shin, Taylor, Yasaman Razeghi, Robert L. Logan, Eric Wallace, and Sameer Singh. 2020. *AutoPrompt: Eliciting Knowledge from Language Models with Automatically Generated Prompts*. <https://doi.org/10.18653/v1/2020.emnlp-main.346>.
- Tan, Qingyu, Hwee Tou Ng, and Lidong Bing. 2023. *Towards Benchmarking and Improving the Temporal Reasoning Capability of Large Language Models*. <https://doi.org/10.18653/v1/2023.acl-long.828>.

- Wei, Jason, Yi Tay, Rishi Bommasani, et al. 2022. “Emergent Abilities of Large Language Models.” *arXiv (Cornell University)*, ahead of print. <https://doi.org/10.48550/arxiv.2206.07682>.
- Xu, Qiantong, Fenglu Hong, B. Li, Changran Hu, Zheng Chen, and Jian Zhang. 2023. “On the Tool Manipulation Capability of Open-Source Large Language Models.” *arXiv.org*, ahead of print. <https://doi.org/10.48550/arXiv.2305.16504>.
- Yan, Lixiang, Lele Sha, Linxuan Zhao, et al. 2023. “Practical and Ethical Challenges of Large Language Models in Education: A Systematic Scoping Review.” *British Journal of Educational Technology*, ahead of print. <https://doi.org/10.1111/bjet.13370>.
- ZHENG, Tianshi, Jiabin Bai, Yicheng Wang, et al. 2024. “CLR-Fact: Evaluating the Complex Logical Reasoning Capability of Large Language Models over Factual Knowledge.” *arXiv.org*, ahead of print. <https://doi.org/10.48550/arXiv.2407.20564>.